

SQL is a Programming Language, Summer 2018

# Handwriting Recognition with SQL (and a tiny bit of web stuff)

Noah Doersing

# *A long time ago in a galaxy very, very close...*

In the early 60s, keyboard proficiency was less widespread, so other input methods were explored.

- 🎥 Alan Kay, "Doing with Images Makes Symbols" (1987):<sup>1</sup>  
**"None of us can type, can you do something about that?"**

---

<sup>1</sup><http://archive.org/details/AlanKeyD1987?start=1439.5>



Instead of implementing an entire interface paradigm, only consider *core problem*: **Recognition of a single-stroke character.**

Slight modifications to some characters to make this work.



Simple D3.js-based web version, connected to a PostgreSQL instance running on my server. Try it on your phones:

[hejnoah.com/handwriting/](https://hejnoah.com/handwriting/)



It yields a JSON-encoded *time series of coordinate pairs*. In order to transfer this into a  $\mathbb{P}$  database...

1. Define variable `pen` via command-line parameter.

```
psql -v pen='[{"x":1, "y":3}, {"x":3, "y":7}]' -f handwriting.sql
```

2. Convert JSON array into tabular representation.

```
WITH RECURSIVE
tablet(pos, x, y) AS (
  SELECT ordinality AS pos, x, y
  FROM  ROWS FROM(jsonb_to_recordset(:'pen') AS (x int, y int)) WITH ORDINALITY
), ...
```

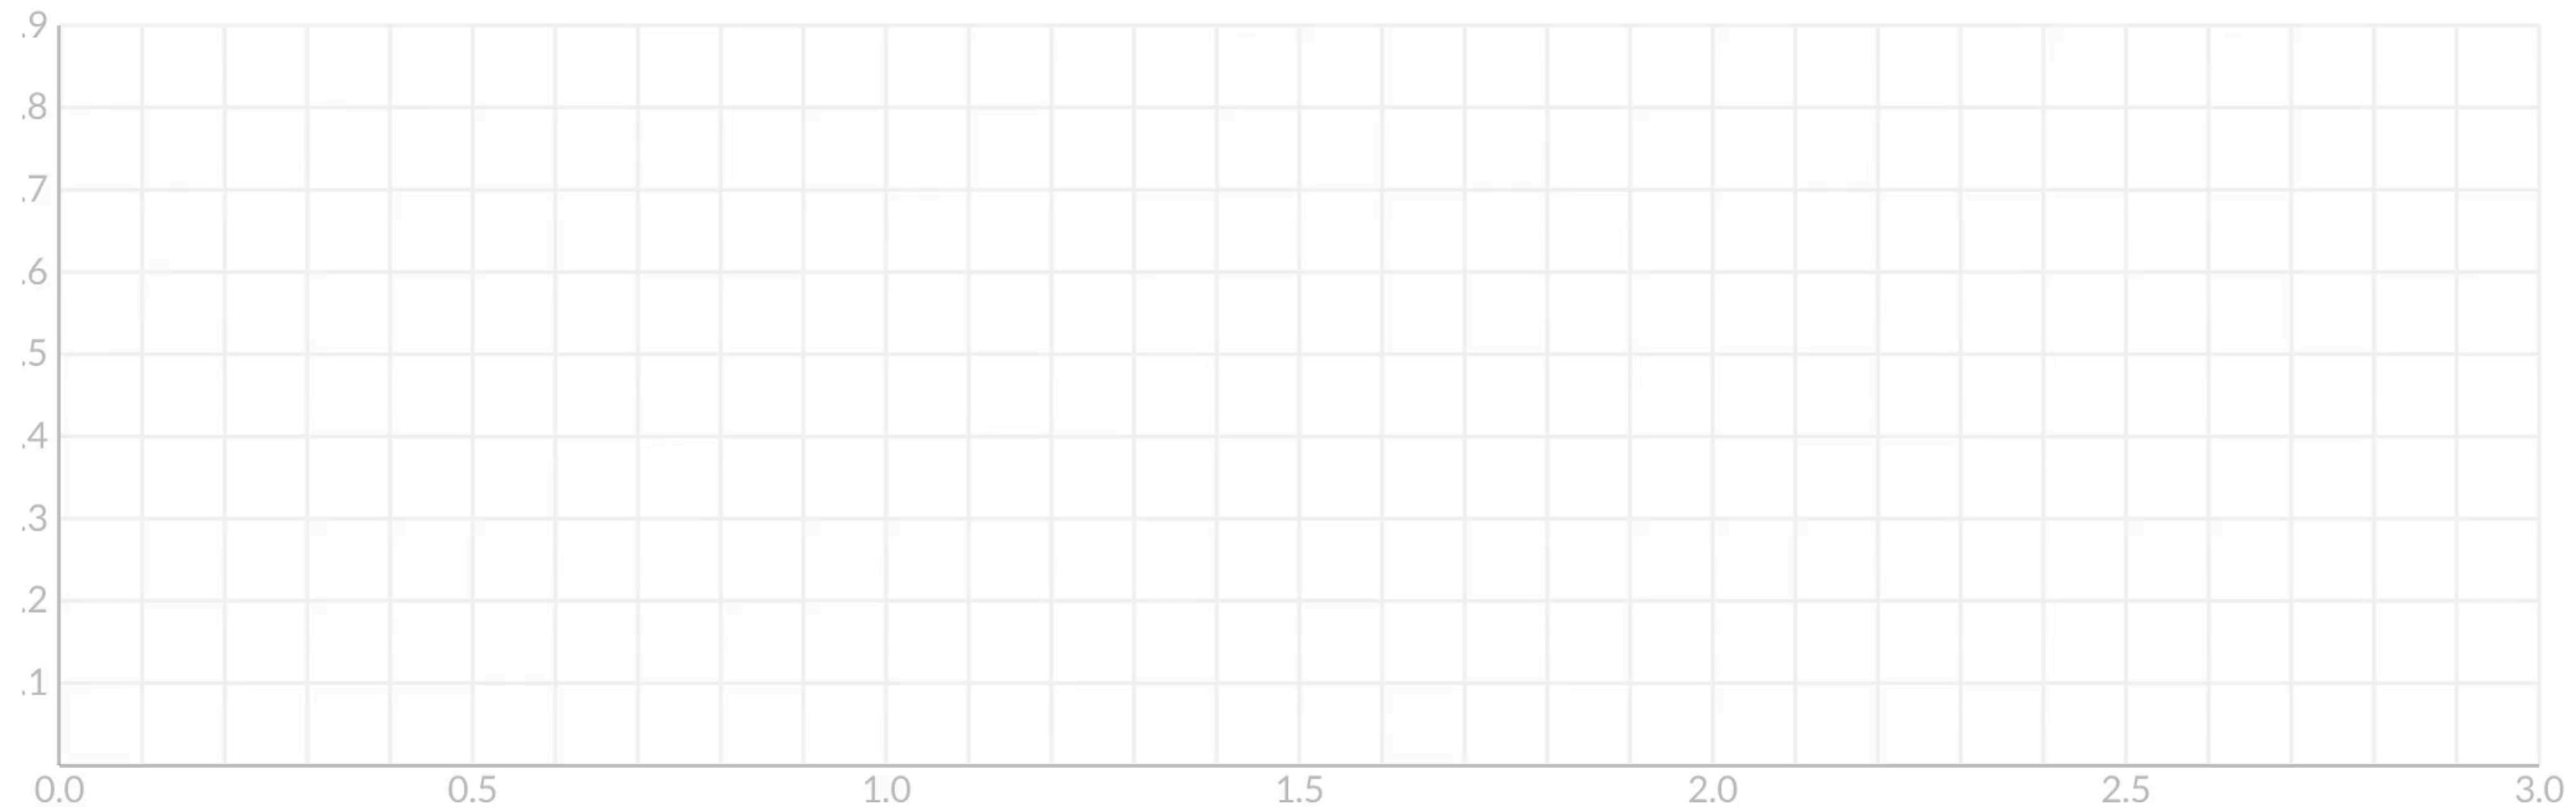
# Approach<sup>2</sup>

1. Smoothing & thinning of stroke
2. Curvature & corner detection
3. Extraction of additional features
4. Descending a hardcoded *decision tree* mapping features to characters

---

<sup>2</sup>[https://www.rand.org/pubs/research\\_memoranda/RM5016.html](https://www.rand.org/pubs/research_memoranda/RM5016.html), <http://jackschaedler.github.io/handwriting-recognition/>

**Smoothing:** Removes *quantization noise*. Compute weighted average of most recently smoothed point and incoming point.



```
smooth(pos, x, y) AS (
    SELECT pos, x :: real, y :: real
    FROM tablet
    WHERE pos = 1

    UNION ALL

    SELECT t.pos,
        (:smoothingfactor * s.x + (1.0 - :smoothingfactor) * t.x) :: real AS x,
        (:smoothingfactor * s.y + (1.0 - :smoothingfactor) * t.y) :: real AS y
    FROM smooth s, tablet t
    WHERE t.pos = s.pos + 1
),
```

Variable **smoothingfactor** must be set before running the query.  
Sensible values: between **0.5** and **0.8**.

```
smooth(pos, x, y) AS (
    SELECT pos, x :: real, y :: real
    FROM tablet
    WHERE pos = 1

    UNION ALL

    SELECT t.pos,
        (:smoothingfactor * s.x + (1.0 - :smoothingfactor) * t.x) :: real AS x,
        (:smoothingfactor * s.y + (1.0 - :smoothingfactor) * t.y) :: real AS y
    FROM smooth s, tablet t
    WHERE t.pos = s.pos + 1
),
```

Variable **smoothingfactor** must be set before running the query.  
Sensible values: between 0.5 and 0.8.

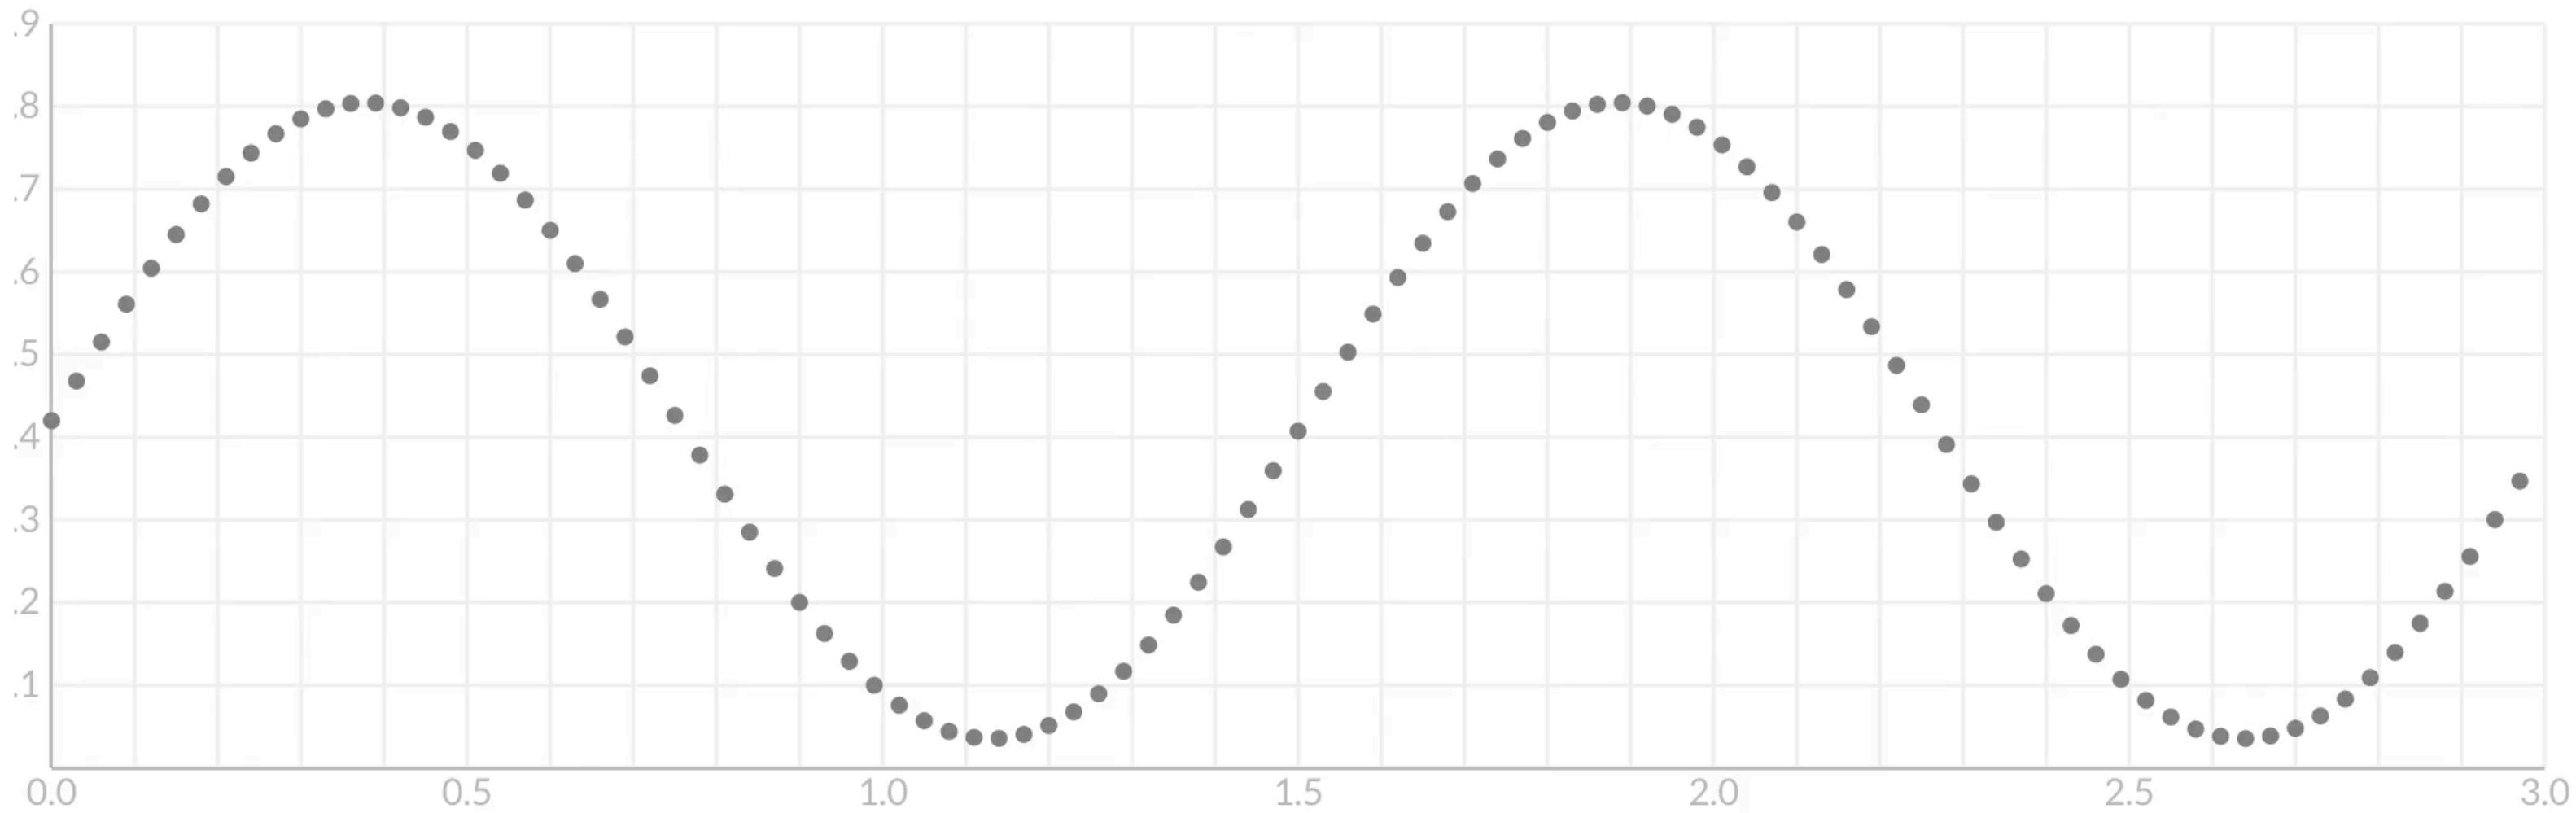
```
smooth(pos, x, y) AS (
    SELECT pos, x :: real, y :: real
    FROM tablet
    WHERE pos = 1

    UNION ALL

    SELECT t.pos,
        (:smoothingfactor * s.x + (1.0 - :smoothingfactor) * t.x) :: real AS x,
        (:smoothingfactor * s.y + (1.0 - :smoothingfactor) * t.y) :: real AS y
    FROM smooth s, tablet t
    WHERE t.pos = s.pos + 1
),
```

Variable **smoothingfactor** must be set before running the query.  
Sensible values: between 0.5 and 0.8.

**Thinning:** Eases further *processing requirements*. Reject points within a certain distance from the most recent accepted point.



```
thin(pos, x, y) AS (
    SELECT *
    FROM   smooth
    WHERE  pos = 1

    UNION ALL

    SELECT *
    FROM  (
        SELECT s.pos, s.x, s.y
        FROM   thin t, smooth s
        WHERE  s.pos > t.pos
        AND    :thinningsize < |/ (s.x - t.x)^2 + (s.y - t.y)^2
        ORDER BY s.pos
        LIMIT 1
    ) AS -
),
```

```
thin(pos, x, y) AS (
    SELECT *
    FROM   smooth
    WHERE  pos = 1

    UNION ALL

    SELECT *
    FROM  (
        SELECT s.pos, s.x, s.y
        FROM   thin t, smooth s
        WHERE  s.pos > t.pos
        AND    :thinningsize < |/ (s.x - t.x)^2 + (s.y - t.y)^2
        ORDER BY s.pos
        LIMIT 1
    ) AS -
),

```

```
thin(pos, x, y) AS (
    SELECT *
    FROM   smooth
    WHERE  pos = 1

    UNION ALL

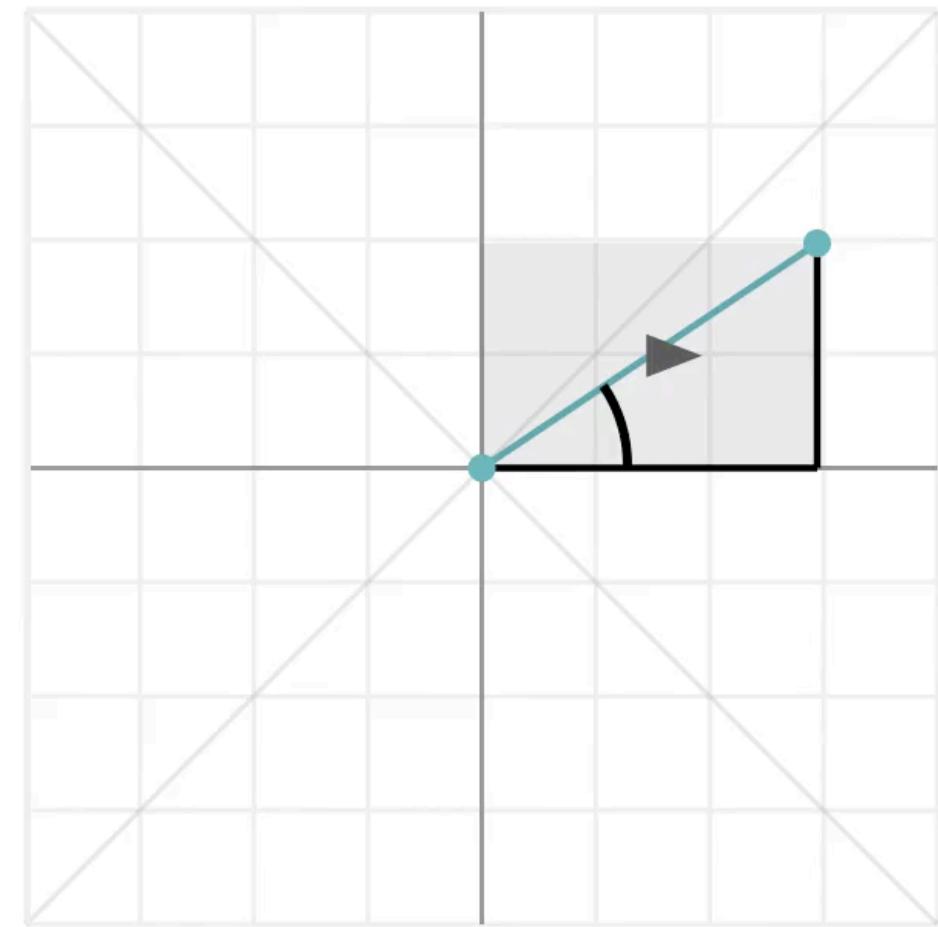
    SELECT *
    FROM  (
        SELECT s.pos, s.x, s.y
        FROM   thin t, smooth s
        WHERE  s.pos > t.pos
        AND    :thinningsize < |/ (s.x - t.x)^2 + (s.y - t.y)^2
        ORDER BY s.pos
        LIMIT 1
    ) AS -
),

```

**Curvature detection:** Compute *cardinal directions* ▲▼◀▶ of line segments between point pairs. Discard sequential duplicates.

Originally done using a set of *inequalities*, but *trigonometry* is more elegant and reasonably fast on modern hardware.

Recall:  $\text{tangent} = \frac{\text{opposite leg}}{\text{adjacent leg}}$



1. Compute angle between every pair of thinned points using window function `lag()`:

```
curve(pos, x, y, direction) AS (
    SELECT pos, x, y,
        degrees(-atan2( y - lag(y) OVER (ORDER BY pos),
                        -x + lag(x) OVER (ORDER BY pos))) + 180
    FROM    thin
),
```

2. Define appropriate **ENUM** type with fancy Unicode triangles:

```
CREATE TYPE cardinal_direction AS ENUM('▶', '▲', '◀', '▼');
```

3. Use **enum\_range()** function (returns the values of its argument's **ENUM** type as an array) to convert **angle / 90** into arrow symbols.

```
cardinal(pos, direction) AS (
    SELECT pos,
        (enum_range(NULL :: cardinal_direction))[(direction / 90) :: int % 4
                                                    + 1]
    FROM   curve
),
```

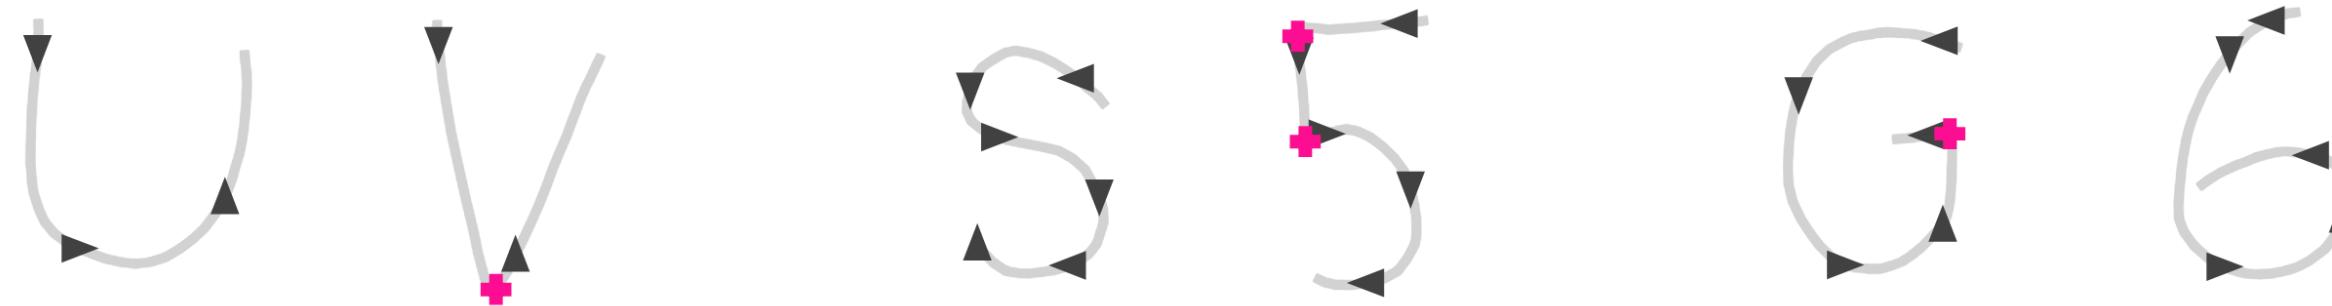
4. Only keep changes of the cardinal direction, i.e. a *new* direction that, additionally, occurs at least *twice in succession*.

```
cardinal_change(pos, direction) AS (
    SELECT pos, direction
    FROM cardinal
    WHERE COALESCE(lag(direction, 2) OVER win <> lag(direction) OVER win,
                    true) -- Prevent NULLs from escaping this term.
        AND lag(direction) OVER win = direction
    WINDOW win AS (ORDER BY pos)
),
```

...except this doesn't work because window functions cannot be used in WHERE clauses. Resort to an awkward subquery.

```
cardinal_change(pos, direction) AS (
    SELECT pos, direction
    FROM  (SELECT pos, direction,
                  COALESCE(lag(direction, 2) OVER win <> lag(direction) OVER win,
                            true) -- Prevent NULLs from escaping this term.
                         AND lag(direction) OVER win = direction
        FROM cardinal
       WINDOW win AS (ORDER BY pos)) AS -(pos, direction, is_new)
  WHERE is_new
),
```

**Corner detection:** Allows discerning between *similar characters*.



Corners lie between two line segments going in the *same direction* and two segments going in a *wildly different direction*, with an optional in-between "turn" segment.



1. Define function that computes angle difference (trivial for angle pairs like  $10^\circ$  and  $20^\circ$ , but less so for  $350^\circ$  and  $30^\circ$ ).

```
CREATE OR REPLACE FUNCTION angdiff(alpha double precision,  
                                beta double precision) RETURNS real AS $$  
SELECT abs(degrees(atan2(sin(radians(alpha - beta)),  
                           cos(radians(alpha - beta))))) :: real;  
--SELECT (180 - abs(abs(alpha - beta) - 180)) :: real;  
$$ LANGUAGE SQL IMMUTABLE;
```

## 2. Detect corners with the help of another awkward subquery.

```
corner(pos, x, y) AS (
  SELECT pos, x, y
  FROM  (SELECT pos, x, y,
                angdiff(lag(direction, 2) OVER win, lag(direction) OVER win) < 22.5
               AND angdiff(lag(direction) OVER win, direction) > :cornerangle
               AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) OR ( -- Immediate direction change OR one-segment turn.
                angdiff(lag(direction, 3) OVER win, lag(direction, 2) OVER win) < 22.5
               AND angdiff(lag(direction, 2) OVER win, direction) > :cornerangle
               AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) AS is_corner
  FROM curve
  WINDOW win AS (ORDER BY pos)) AS -(pos, x, y, is_corner)
 WHERE is_corner
),
```

## 2. Detect corners with the help of another awkward subquery.

```
corner(pos, x, y) AS (
  SELECT pos, x, y
  FROM  (SELECT pos, x, y,
                angdiff(lag(direction, 2) OVER win, lag(direction) OVER win) < 22.5
               AND angdiff(lag(direction) OVER win, direction) > :cornerangle
               AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) OR ( -- Immediate direction change OR one-segment turn.
                angdiff(lag(direction, 3) OVER win, lag(direction, 2) OVER win) < 22.5
               AND angdiff(lag(direction, 2) OVER win, direction) > :cornerangle
               AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) AS is_corner
  FROM curve
  WINDOW win AS (ORDER BY pos)) AS -(pos, x, y, is_corner)
 WHERE is_corner
),
```

## 2. Detect corners with the help of another awkward subquery.

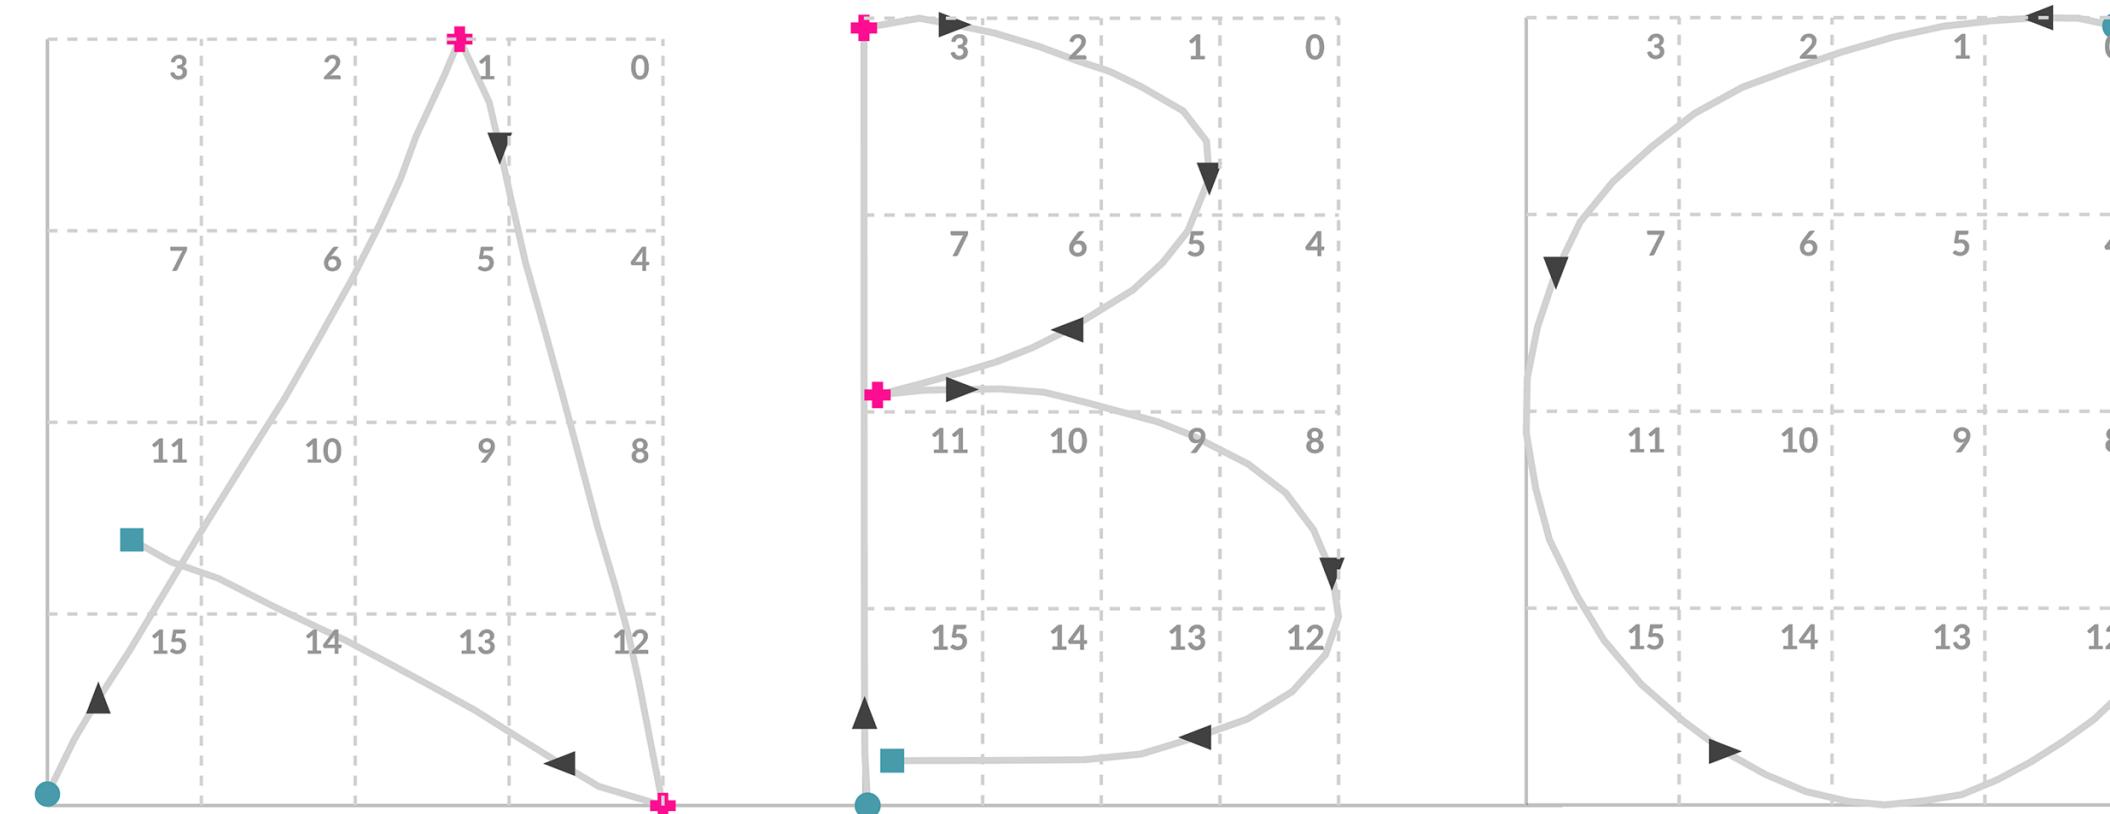
```
corner(pos, x, y) AS (
  SELECT pos, x, y
  FROM  (SELECT pos, x, y,
                angdiff(lag(direction, 2) OVER win, lag(direction) OVER win) < 22.5
              AND angdiff(lag(direction) OVER win, direction) > :cornerangle
              AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) OR ( -- Immediate direction change OR one-segment turn.
                angdiff(lag(direction, 3) OVER win, lag(direction, 2) OVER win) < 22.5
              AND angdiff(lag(direction, 2) OVER win, direction) > :cornerangle
              AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) AS is_corner
  FROM curve
  WINDOW win AS (ORDER BY pos)) AS -(pos, x, y, is_corner)
 WHERE is_corner
),
```

## 2. Detect corners with the help of another awkward subquery.

```
corner(pos, x, y) AS (
  SELECT pos, x, y
  FROM  (SELECT pos, x, y,
                angdiff(lag(direction, 2) OVER win, lag(direction) OVER win) < 22.5
               AND angdiff(lag(direction) OVER win, direction) > :cornerangle
               AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) OR ( -- Immediate direction change OR one-segment turn.
                angdiff(lag(direction, 3) OVER win, lag(direction, 2) OVER win) < 22.5
               AND angdiff(lag(direction, 2) OVER win, direction) > :cornerangle
               AND angdiff(direction, lead(direction) OVER win) < 22.5
            ) AS is_corner
  FROM curve
  WINDOW win AS (ORDER BY pos)) AS -(pos, x, y, is_corner)
 WHERE is_corner
),
```

**Feature extraction:** Extract some *supplementary features* that will help discern between different characters later on: Start point, end point, aspect ratio, ...

Transform from absolute pixel positions to 4x4 grid segments.



1. Define *axis-aligned bounding box* around pen stroke and gather some statistics.

```
aabb(xmin, xmax, ymin, ymax, aspect, width, height, centerx, centery) AS (
  SELECT min(x),
         max(x),
         min(y),
         max(y),
         (max(y) - min(y)) / greatest(1, (max(x) - min(x))), -- Prevent n/0.
         max(x) - min(x),
         max(y) - min(y),
         min(x) + (max(x) - min(x)) / 2,
         min(y) + (max(y) - min(y)) / 2
  FROM   smooth
),
```

2. Define function for *transforming* coordinates of start, end and corners into grid positions.

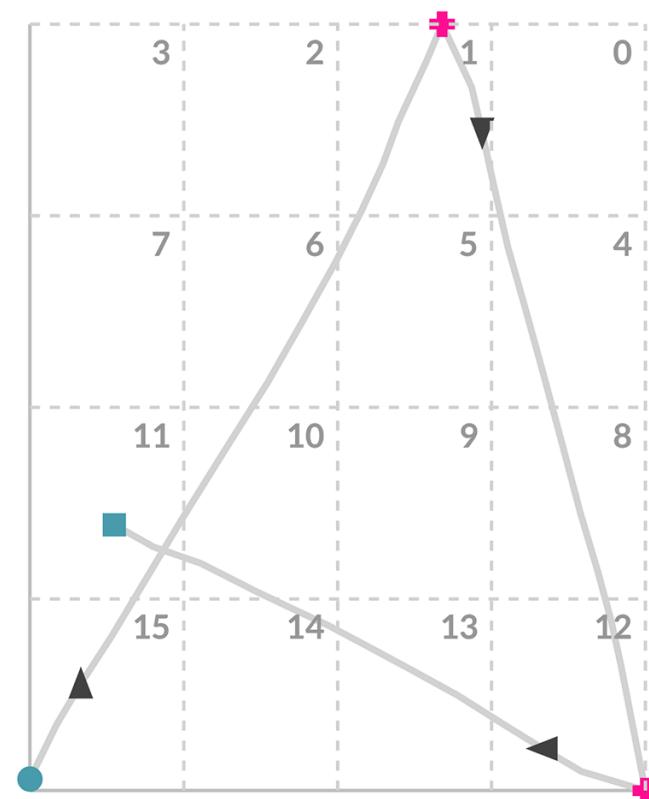
```
CREATE OR REPLACE FUNCTION gridpos(width real, height real,  
                                  xmin real, ymin real,  
                                  x real, y real) RETURNS int AS $$  
SELECT greatest(0,  
                15 - (  floor(4 * (x - xmin) / (width + 1)) :: int)  
                      + 4 * (floor(4 * (y - ymin) / (height + 1)) :: int)));  
$$ LANGUAGE SQL IMMUTABLE;
```

3. Apply this function and thus create CTEs `start_grid`, `stop_grid` and `corner_grid`. Not shown.

4. Accumulate extracted features. This single-row table contains all information required to identify the drawn character.

```
features(directions, start, stop, corners, width, height, aspect, center) AS (
    SELECT (SELECT array_agg(direction ORDER BY pos)
        FROM cardinal_change),
    (TABLE start_grid),
    (TABLE stop_grid),
    (SELECT COALESCE(array_agg(n ORDER BY pos), '{}')
        FROM corner_grid),
    width,
    height,
    aspect,
    point(centerx, centery)
    FROM aabb
),
```

directions	start	stop	corners	width	height	aspect	center
$\{\uparrow, \downarrow, \leftarrow\}$	15	11	{1,12}	36.7595	66.1965	1.8008	(55.3797, 64.0982)



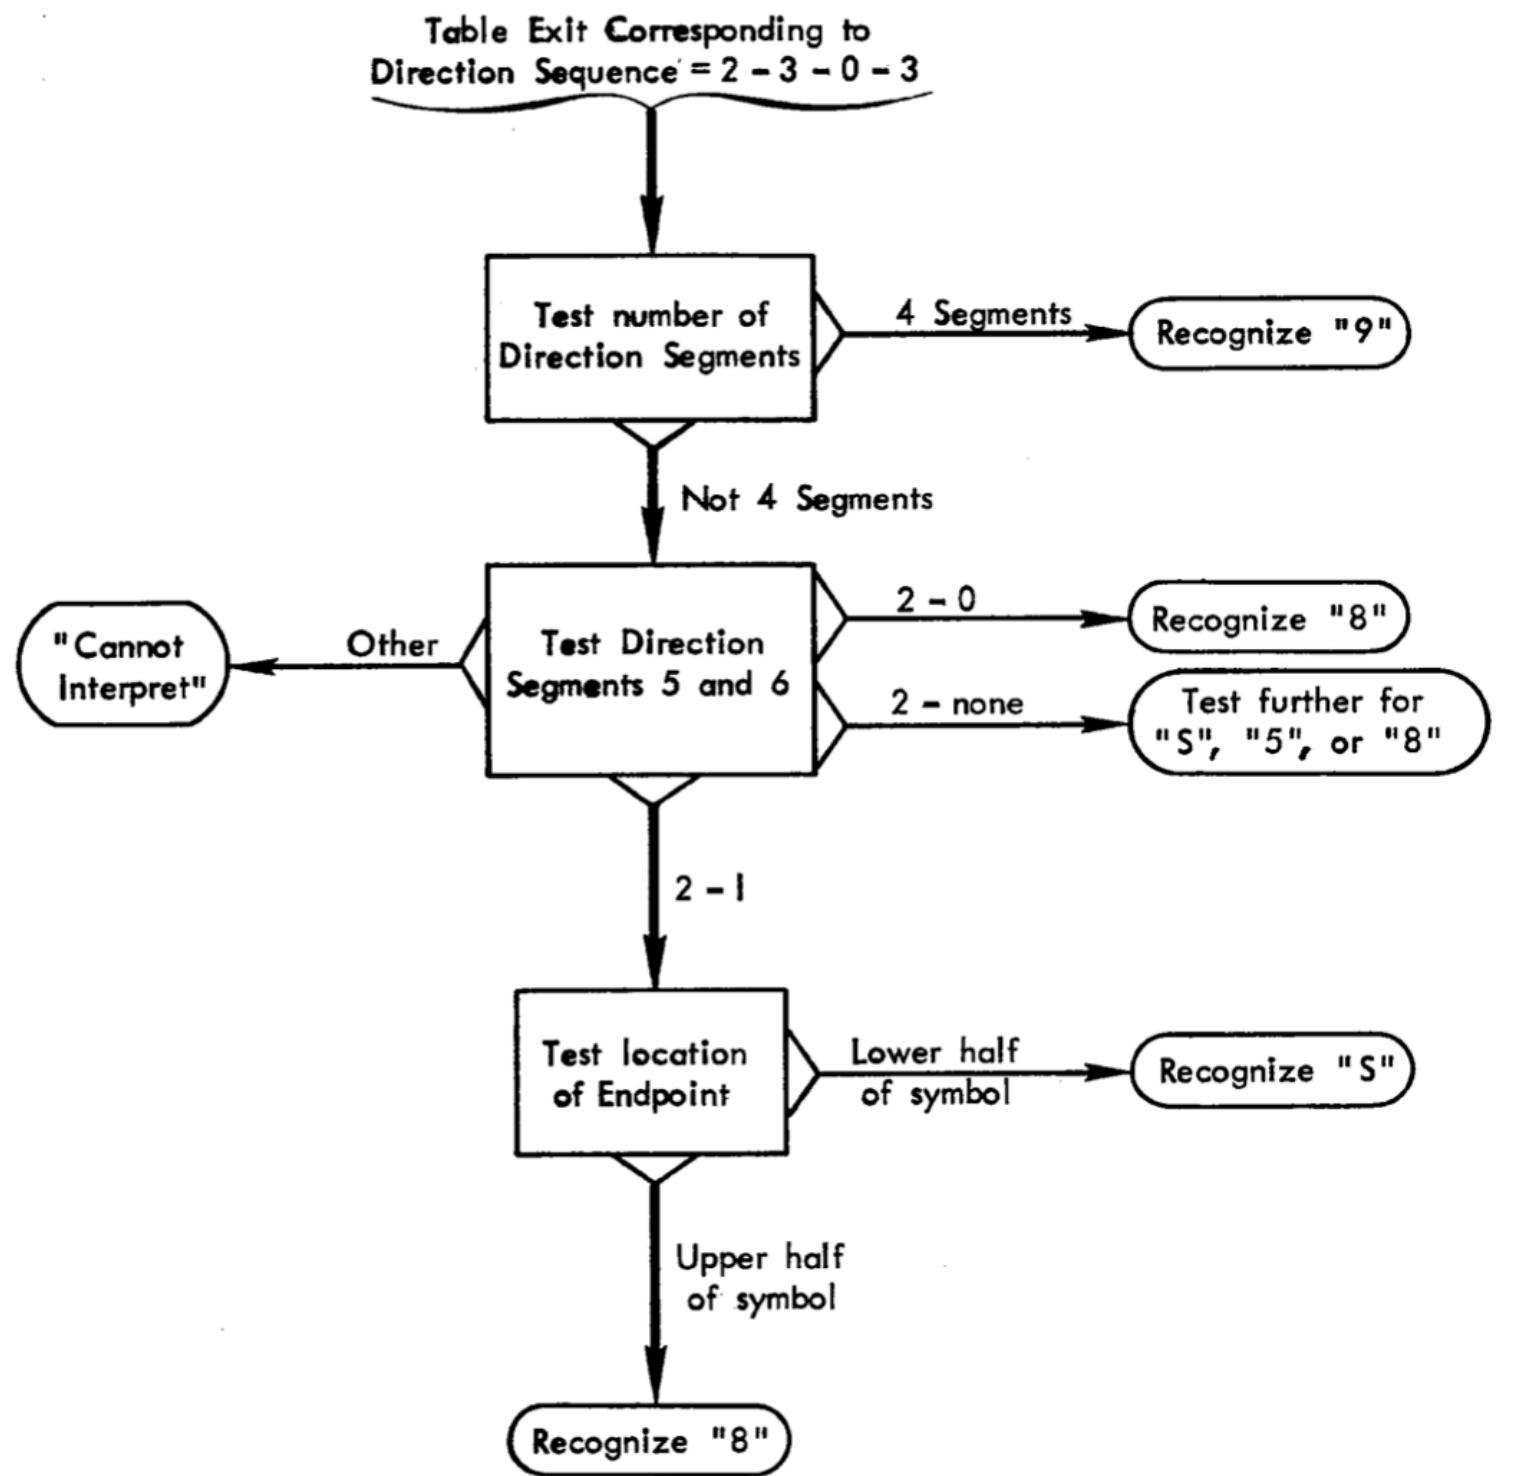


FIG. 6 - Recognition of the Track of Figure 2.

**Decision tree:** Depending on the first four cardinal directions, descend custom-tailored subtree.

Incredibly *awkward* in SQL.  
Imagine 1000+ lines of deeply nested and probably buggy CASE WHENs.

A better approach is to do this in *two discrete stages*:

1. Determine **candidate characters** based on the pen stroke's *first four cardinal directions*.
2. Settle on **best fit** by matching *selected stroke features*.

Less flexible than the procedure presented in the original memo, but significantly more *idiomatic* and concise.

# 1. Determine initial set of potential characters.

```
CREATE TABLE lookup_candidates (
    first_four_directions cardinal_direction[],
    candidate_characters char[])
);

INSERT INTO lookup_candidates VALUES
('{"▼"}', ' {"I"}'),
('{"▼", "◀"}', ' {"J"}'),
('{"▼", "◀", "▲"}', ' {"O", "J", "X", "U"}'),
('{"▼", "◀", "▲", "▶"}', ' {"X", "O", "U"}'),
('{"▼", "◀", "▶", "▲"}', ' {"X"}'),
('{"▼", "▶"}', ' {"L"}'),
('{"▼", "▶", "◀"}', ' {"6"}'),
('{"▼", "▶", "◀", "▼"}', ' {"4"}'),
('{"▼", "▶", "▲"}', ' {"O", "U"}'),
...

```

2. Find best fit by feature matching. Read **NULLs** as "don't care".

```
CREATE TABLE lookup_bestfit (
    candidate_characters char[], 
    character            char,
    start                int,
    stop                 int,
    corners              int[],
    aspect_range         numrange
);
INSERT INTO lookup_bestfit -- All single-character patterns from first lookup.
...
INSERT INTO lookup_bestfit VALUES
('{"M","N"}', 'M', NULL, 12, NULL, NULL), -- End point bottom right.
('{"M","N"}', 'N', NULL, 0,  NULL, NULL), -- End point top right.
...
```

3. Tie it all together: Consult *both* lookup tables in order.

```
character(character) AS (
  SELECT l.character
  FROM   features f, lookup_candidates c, lookup_bestfit l
  WHERE  f.directions[1:4] = c.first_four_directions
  AND    c.candidate_characters = l.candidate_characters
  AND    (l.start IS NULL          OR f.start = l.start)
  AND    (l.stop  IS NULL          OR f.stop  = l.stop)
  AND    (l.corners IS NULL        OR f.corners = l.corners)
  AND    (l.aspect_range IS NULL    OR l.aspect_range @> f.aspect :: numeric)
),
```

### 3. Tie it all together: Consult *both* lookup tables in order.

```
character(character) AS (
  SELECT l.character
  FROM   features f, lookup_candidates c, lookup_bestfit l
  WHERE  f.directions[1:4] = c.first_four_directions
  AND    c.candidate_characters = l.candidate_characters
  AND    (l.start IS NULL          OR f.start = l.start)
  AND    (l.stop IS NULL           OR f.stop = l.stop)
  AND    (l.corners IS NULL        OR f.corners = l.corners)
  AND    (l.aspect_range IS NULL   OR l.aspect_range @> f.aspect :: numeric)
),
```

### 3. Tie it all together: Consult *both* lookup tables in order.

```
character(character) AS (
  SELECT l.character
  FROM   features f, lookup_candidates c, lookup_bestfit l
  WHERE  f.directions[1:4] = c.first_four_directions
  AND    c.candidate_characters = l.candidate_characters
  AND    (l.start IS NULL           OR f.start = l.start)
  AND    (l.stop  IS NULL           OR f.stop  = l.stop)
  AND    (l.corners IS NULL         OR f.corners = l.corners)
  AND    (l.aspect_range IS NULL    OR l.aspect_range @> f.aspect :: numeric)
),
```

### 3. Tie it all together: Consult *both* lookup tables in order.

```
character(character) AS (
  SELECT l.character
  FROM   features f, lookup_candidates c, lookup_bestfit l
  WHERE  f.directions[1:4] = c.first_four_directions
  AND    c.candidate_characters = l.candidate_characters
  AND    (l.start IS NULL          OR f.start = l.start)
  AND    (l.stop  IS NULL          OR f.stop  = l.stop)
  AND    (l.corners IS NULL        OR f.corners = l.corners)
  AND    (l.aspect_range IS NULL    OR l.aspect_range @> f.aspect :: numeric)
),
```

### 3. Tie it all together: Consult *both* lookup tables in order.

```
character(character) AS (
  SELECT l.character
  FROM   features f, lookup_candidates c, lookup_bestfit l
  WHERE  f.directions[1:4] = c.first_four_directions
  AND    c.candidate_characters = l.candidate_characters
  AND    (l.start IS NULL          OR f.start = l.start)
  AND    (l.stop IS NULL           OR f.stop = l.stop)
  AND    (l.corners IS NULL        OR f.corners = l.corners)
  AND    (l.aspect_range IS NULL   OR l.aspect_range @> f.aspect :: numeric)
),
```

# Performance: Most algorithms designed in the 60's will run on modern hardware just *fine* – this one is no exception:

The whole chain of queries terminates after  $\sim 15$  ms.

## SQL is, *after all*, a programming language.

```

QUERY PLAN

CTE Scan on debug (cost=601.52..601.60 rows=4 width=172) (actual time=10.615..10.705 rows=1 loops=1)
  CTE tablet
    -> Function Scan on jsonb_to_recordset (cost=0.00..1.00 rows=100 width=16) (actual time=0.104..0.117 rows=64 loops=1)
  CTE smooth
    -> Recursive Union (cost=0.00..36.53 rows=101 width=16) (actual time=0.115..2.662 rows=64 loops=1)
      -> CTE Scan on tablet_t (cost=0.00..2.25 rows=1 width=16) (actual time=0.114..0.161 rows=1 loops=1)
        Filter: (pos < 1)
        Rows Removed by Filter: 63
      -> Hash Join (cost=0.33..3.23 rows=10 width=16) (actual time=0.025..0.037 rows=1 loops=64)
        Hash Cond: (t.pos = (s_1.pos + 1))
          -> CTE Scan on tablet_t (cost=0.00..2.00 rows=100 width=16) (actual time=0.000..0.011 rows=64 loops=64)
          -> Hash (cost=0.20..0.20 rows=10 width=16) (actual time=0.002..0.002 rows=1 loops=64)
            Buckets: 1024 Batches: 1 Memory Usage: 9KB
          -> WorkTable Scan on smooth_s (cost=0.00..0.20 rows=10 width=16) (actual time=0.000..0.001 rows=1 loops=64)
  CTE thin
    -> Recursive Union (cost=0.00..474.82 rows=11 width=16) (actual time=0.003..2.681 rows=23 loops=1)
      -> CTE Scan on smooth (cost=0.00..2.27 rows=1 width=16) (actual time=0.002..0.021 rows=1 loops=1)
        Filter: (pos = 1)
        Rows Removed by Filter: 63
      -> Limit (cost=47.22..47.22 rows=1 width=16) (actual time=0.113..0.113 rows=1 loops=23)
        -> Sort (cost=47.22..47.50 rows=12 width=16) (actual time=0.109..0.109 rows=1 loops=23)
          Sort Key: s_1.pos
          Sort Method: quicksort Memory: 25KB
        -> Nested Loop (cost=0.00..46.66 rows=12 width=16) (actual time=0.017..0.084 rows=39 loops=23)
          Join Filter: ((s_1.pos > t_1.pos) AND ('5'::double precision < (/ (((s_1.x - t_1.x)::double precision ^ '2'::double precision) + ((s_1.y - t_1.y)::double precision ^ '2'::double precision)))))
          Rows Removed by Join Filter: 25
        -> WorkTable Scan on thin_t_1 (cost=0.00..2.02 rows=101 width=16) (actual time=0.000..0.016 rows=64 loops=23)
        -> CTE Scan on smooth_s_1 (cost=0.00..2.02 rows=101 width=16) (actual time=0.000..0.016 rows=64 loops=23)
  CTE curve
    -> WindowAgg (cost=0.41..0.88 rows=11 width=24) (actual time=2.846..3.071 rows=23 loops=1)
      -> Sort (cost=0.41..0.44 rows=11 width=16) (actual time=2.797..2.810 rows=23 loops=1)
        Sort Key: thin_pos
        Sort Method: quicksort Memory: 25KB
      -> CTE Scan on thin (cost=0.00..2.22 rows=11 width=16) (actual time=0.004..2.702 rows=23 loops=1)
  CTE cardinal
    -> CTE Scan on curve (cost=0.00..0.36 rows=11 width=12) (actual time=3.256..3.384 rows=23 loops=1)
  CTE cardinal_change
    -> Subquery Scan on _ (cost=0.41..0.80 rows=6 width=12) (actual time=3.959..4.054 rows=3 loops=1)
      Filter: _1_is_new
      Rows Removed by Filter: 20
    -> WindowAgg (cost=0.41..0.69 rows=11 width=13) (actual time=3.949..4.045 rows=23 loops=1)
      -> Sort (cost=0.41..0.44 rows=11 width=12) (actual time=3.929..3.933 rows=23 loops=1)
        Sort Key: cardinal_pos
        Sort Method: quicksort Memory: 25KB
      -> CTE Scan on cardinal (cost=0.00..0.22 rows=11 width=12) (actual time=3.260..3.384 rows=23 loops=1)
  CTE corner
    -> Subquery Scan on _1 (cost=0.41..2.61 rows=6 width=16) (actual time=0.451..0.566 rows=2 loops=1)
      Filter: _1_is_corner
      Rows Removed by Filter: 21
    -> WindowAgg (cost=0.41..2.58 rows=11 width=17) (actual time=0.287..0.554 rows=23 loops=1)
      -> Sort (cost=0.41..0.44 rows=11 width=24) (actual time=0.182..0.186 rows=23 loops=1)
        Sort Key: curve_1_pos
        Sort Method: quicksort Memory: 25KB
      -> CTE Scan on curve (cost=0.00..0.22 rows=11 width=16) (actual time=0.001..0.055 rows=23 loops=1)
  CTE aabb
    -> Aggregate (cost=6.57..6.61 rows=1 width=44) (actual time=2.792..2.792 rows=1 loops=1)
    -> CTE Scan on smooth_smooth_1 (cost=0.00..2.02 rows=101 width=8) (actual time=0.116..2.787 rows=64 loops=1)
  CTE start_grid
    -> Limit (cost=3.56..3.61 rows=1 width=12) (actual time=0.113..0.114 rows=1 loops=1)
      -> Result (cost=3.56..3.88 rows=101 width=12) (actual time=0.106..0.106 rows=1 loops=1)
        -> Sort (cost=3.56..3.81 rows=101 width=32) (actual time=0.097..0.097 rows=1 loops=1)
          Sort Key: s_2_pos
          Sort Method: top-N heapsort Memory: 25KB
        -> Nested Loop (cost=0.00..3.05 rows=101 width=22) (actual time=0.007..0.053 rows=64 loops=1)
          -> CTE Scan on aabb_a (cost=0.00..0.02 rows=1 width=16) (actual time=0.001..0.002 rows=1 loops=1)
          -> CTE Scan on smooth_s_2 (cost=0.00..2.02 rows=101 width=16) (actual time=0.001..0.017 rows=64 loops=1)
  CTE stop_grid
    -> Limit (cost=3.56..3.61 rows=1 width=12) (actual time=0.190..0.191 rows=1 loops=1)
      -> Result (cost=3.56..3.88 rows=101 width=12) (actual time=0.190..0.190 rows=1 loops=1)
        -> Sort (cost=3.56..3.81 rows=101 width=32) (actual time=0.185..0.185 rows=1 loops=1)
          Sort Key: s_3_pos
          Sort Method: top-N heapsort Memory: 25KB
        -> Nested Loop (cost=0.00..3.05 rows=101 width=32) (actual time=0.002..0.105 rows=64 loops=1)
          -> CTE Scan on aabb_a_1 (cost=0.00..0.02 rows=1 width=16) (actual time=0.001..0.001 rows=1 loops=1)
          -> CTE Scan on smooth_s_3 (cost=0.00..2.02 rows=101 width=16) (actual time=0.000..0.022 rows=64 loops=1)
  CTE corner_grid
    -> Unique (cost=0.52..0.55 rows=1 width=12) (actual time=0.083..0.087 rows=2 loops=1)
      -> Sort (cost=0.52..0.53 rows=1 width=12) (actual time=0.082..0.084 rows=2 loops=1)
        Sort Key: (GREATEST((15 - ((c.x - a_2.xmin) * (a_2.width + '1'::double precision)))::integer + (4 * (floor((('4'::double precision * (c.y - a_2.ymin)) / (a_2.height + '1'::double precision))))::integer
        Sort Method: quicksort Memory: 25KB
      -> Nested Loop (cost=0.00..0.44 rows=1 width=12) (actual time=0.473..0.593 rows=2 loops=1)
        -> CTE Scan on aabb_a_2 (cost=0.00..0.02 rows=1 width=16) (actual time=0.003..0.004 rows=1 loops=1)
        -> CTE Scan on corner_c (cost=0.00..0.12 rows=6 width=16) (actual time=0.454..0.571 rows=2 loops=1)
  CTE features
    -> CTE Scan on aabb (cost=0.34..0.36 rows=1 width=100) (actual time=8.370..8.371 rows=1 loops=1)
    InitPlan 12 (returns $13)
      -> Aggregate (cost=14..0.15 rows=1 width=32) (actual time=4.150..4.150 rows=1 loops=1)
        -> CTE Scan on cardinal_change (cost=0.00..0.12 rows=6 width=12) (actual time=3.963..4.062 rows=3 loops=1)
    InitPlan 13 (returns $14)
      -> CTE Scan on start_grid (cost=0.00..0.02 rows=1 width=4) (actual time=0.148..0.149 rows=1 loops=1)
    InitPlan 14 (returns $15)
      -> CTE Scan on stop_grid (cost=0.00..0.02 rows=1 width=4) (actual time=0.193..0.194 rows=1 loops=1)
    InitPlan 15 (returns $16)
      -> Aggregate (cost=0.14..0.15 rows=1 width=32) (actual time=0.952..0.953 rows=1 loops=1)
        -> CTE Scan on corner_grid (cost=0.00..0.12 rows=6 width=12) (actual time=0.865..0.813 rows=2 loops=1)
  CTE character
    -> Hash Join (cost=22.22..40.22 rows=1 width=8) (actual time=0.560..1.346 rows=1 loops=1)
      Hash Cond: (l.candidate_characters = c_1.candidate_characters)
      Join Filter: ((l.start IS NULL) OR (f.start = l.start) AND ((l.stop IS NULL) OR (f.stop = l.stop)) AND ((l.corners IS NULL) OR (f.corners = l.corners)) AND ((l.last_direction IS NULL) OR (f.directions[array_length(f.directions, 1)] = l.last_
      -> Seq Scan on lookup_bestfit_l (cost=0.00..15.60 rows=568 width=16) (actual time=0.096..0.389 rows=353 loops=1)
      -> Hash (cost=22.17..22.17 rows=4 width=108) (actual time=0.270..0.270 rows=1 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 9KB
        -> Hash Join (cost=0.03..0.23 rows=1 width=108) (actual time=0.209..0.230 rows=1 loops=1)
          Hash Cond: (l.candidate_characters = c_1.candidate_characters)
          Join Filter: ((l.start IS NULL) OR (f.start = l.start) AND ((l.stop IS NULL) OR (f.stop = l.stop)) AND ((l.corners IS NULL) OR (f.corners = l.corners)) AND ((l.last_direction IS NULL) OR (f.directions[array_length(f.directions, 1)] = l.last_
            -> Seq Scan on lookup_candidate_c_1 (cost=0.00..18.80 rows=880 width=64) (actual time=0.033..0.057 rows=70 loops=1)
            -> Hash (cost=0.02..0.02 rows=1 width=108) (actual time=0.034..0.034 rows=1 loops=1)
              Buckets: 1024 Batches: 1 Memory Usage: 9KB
              -> CTE Scan on features_f (cost=0.00..0.02 rows=1 width=76) (actual time=0.006..0.007 rows=1 loops=1)
  CTE debug
    -> Hash Join (cost=0.07..29.58 rows=4 width=172) (actual time=10.580..10.669 rows=1 loops=1)
      Hash Cond: (l.debug_candidates.first_four_directions = features.directions[1:4])
      Join Filter: (l.debug_candidates.first_four_directions = features.directions[1:4])
      -> CTE Scan on "character" (cost=0.00..0.02 rows=1 width=8) (actual time=0.568..1.1356 rows=1 loops=1)
    InitPlan 20 (returns $21)
      -> CTE Scan on "character" character_1 (cost=0.00..0.02 rows=1 width=8) (actual time=0.001..0.001 rows=1 loops=1)
      -> Seq Scan on lookup_candidate (cost=0.00..18.80 rows=880 width=64) (actual time=0.022..0.044 rows=70 loops=1)
      -> Hash (cost=0.02..0.02 rows=1 width=108) (actual time=0.047..0.474 rows=1 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 9KB
        -> CTE Scan on features_f (cost=0.00..0.02 rows=1 width=108) (actual time=0.388..0.382 rows=1 loops=1)
  SubPlan 1 (returns $22)
    -> Aggregate (cost=1.75..1.76 rows=1 width=32) (actual time=0.278..0.279 rows=1 loops=1)
      -> Function Scan on unnest (cost=0.00..1.00 rows=100 width=40) (actual time=0.183..0.184 rows=1 loops=1)
Planning time: 5.633 ms
Execution time: 12.077 ms

```

# Demo

Handwriting Recognition with SQL  
Noah Doersing